

Exploratory Data Analysis with R

Data Manipulation Using Dplyr

Xuemao Zhang
East Stroudsburg University


September 30, 2022

Outline

- Overview
- Filter observations (rows)
 - ▶ Removing missing values
- Select/drop variables (columns)
- The `%>%` operator
- Add/remove variables
- Rename variables
- Sort data
- Remove duplicate rows

Overview

- dplyr, part of **Tidyverse** is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges.
- The dplyr package is one of the most powerful and popular package in R.
- The author, Hadley Wickham, is also the author of the package 'ggplot2'.
- R for Data Science (O'Reilly 2017) by Hadley Wickham
- Free Online: <https://r4ds.had.co.nz/>



Overview

- Data Exploration = Data manipulation + Data visualization.
- **Data wrangling** (a term used in data science) is the process of transforming/mapping data from raw format into ready-to-analyze format.
- Besides `ggplot2()` for data visualization, Hadley Wickham created a series of R packages for data wrangling, including
 - ▶ `tidyr` for reshaping your data for plotting and use by different R functions
 - ▶ `tibble` for better ways to create, print and subset data frames
 - ▶ `dplyr` for data manipulation will be covered in this course
- The `dplyr` package also interfaces well with tibbles.
- `dplyr` functions process **faster** than base R functions. It is because `dplyr` functions were written in a computationally efficient manner. They are also more stable in the syntax and better supports data frames than vectors.

Overview

- There are 8 fundamental data manipulation verbs in `dplyr` that you will use to do most of your data manipulations.
 - ▶ `mutate()` and `transmute()`: Add/create new variables.
 - ▶ `select()`: Select columns (variables) by their names.
 - ▶ `filter()`: Pick rows (observations/samples) based on their values.
 - ▶ `distinct()`: Remove duplicate rows.
 - ▶ `arrange()`: Reorder the rows.
 - ▶ `rename()`: Rename columns.
 - ▶ `summarise()`: Compute statistical summaries (e.g., computing the mean or the sum) and thus reduce multiple values down to a single summary. It is similar to `R::base::aggregate`. - `group_by()` to group variables for `summarise`

Recall. The double colon `::` is used to specify the package where a function is from: `packagename::functionname()`.

Overview

- All these functions work similarly as follow:
 - ▶ The first argument is a data frame
 - ▶ The subsequent arguments are comma separated list of unquoted variable names and the specification of what you want to do
 - ▶ The result is a new data frame
- A special feature in dplyr is that you can chain your data manipulation operations using the pipe operator (`%>%`).

Note. dplyr package allows to use the forward-pipe chaining operator (`%>%`) for combining multiple operations. For example, `x %>% f` is equivalent to `f(x)`. Using the pipe (`%>%`), the output of each operation is passed to the next operation. This makes R programming easy:

# input	+-----+		+-----+		+-----+		result
# data	%>%	verb	%>%	verb	%>%	verb	-> data
# frame	+-----+		+-----+		+-----+		frame

Creating a data.frame to work with

Again, here we use one of the datasets that comes with R called `mtcars` create a toy data.frame named `df`

```
data(mtcars);  
dfm = mtcars; # to save original
```

No rownames in tibbles!

In the “tidy” data format, all information of interest is a variable (not a name). **as of tibble 2.0, rownames are removed.** For example, `mtcars` has each car name as a row name:

```
head(dfm, 2)
```

```
##           mpg cyl  disp  hp  drat    wt  qsec vs  am  gear carb
## Mazda RX4    21   6  160 110   3.9 2.620 16.46  0   1    4    4
## Mazda RX4 Wag 21   6  160 110   3.9 2.875 17.02  0   1    4    4
```

```
library(tibble)
head(as_tibble(dfm), 2)
```

```
## # A tibble: 2 x 11
##   mpg     cyl  disp    hp  drat    wt  qsec    vs    am  gear carb
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1    21     6   160   110   3.9  2.62  16.5     0     1     4     4
## 2    21     6   160   110   3.9  2.88  17.0     0     1     4     4
```


No rownames in tibbles!

If you run into this, use `rownames_to_column` to add it before turning it into a tibble to keep them:

```
dfm = rownames_to_column(dfm, var = "car");  
dfm = as_tibble(dfm);  
dfm;
```

```
## # A tibble: 32 x 12
```

	car	mpg	cyl	disp	hp	drat	wt	qsec	vs	<
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<
## 1	Mazda RX4	21	6	160	110	3.9	2.62	16.5	0	
## 2	Mazda RX4 ~	21	6	160	110	3.9	2.88	17.0	0	
## 3	Datsun 710	22.8	4	108	93	3.85	2.32	18.6	1	
## 4	Hornet 4 D~	21.4	6	258	110	3.08	3.22	19.4	1	
## 5	Hornet Spo~	18.7	8	360	175	3.15	3.44	17.0	0	
## 6	Valiant	18.1	6	225	105	2.76	3.46	20.2	1	
## 7	Duster 360	14.3	8	360	245	3.21	3.57	15.8	0	
## 8	Merc 240D	24.4	4	147.	62	3.69	3.19	20	1	
## 9	Merc 230	22.8	4	141.	95	3.92	3.15	22.9	1	
## 10	Merc 280	19.2	6	168.	123	3.92	3.44	18.3	1	

```
## # ... with 22 more rows
```

No rownames in tibbles!

- The function `dplyr::glimpse` is a little like `str()` applied to a data frame but it tries to show you as much data as possible.

```
dplyr::glimpse(dfm);
```

```
## Rows: 32
## Columns: 12
## $ car   <chr> "Mazda RX4", "Mazda RX4 Wag", "Datsun 710", "Hornet
## $ mpg   <dbl> 21.0, 21.0, 22.8, 21.4, 18.7, 18.1, 14.3, 24.4, 22.8
## $ cyl   <dbl> 6, 6, 4, 6, 8, 6, 8, 4, 4, 6, 6, 8, 8, 8, 8, 8, 4
## $ disp  <dbl> 160.0, 160.0, 108.0, 258.0, 360.0, 225.0, 360.0, 146
## $ hp    <dbl> 110, 110, 93, 110, 175, 105, 245, 62, 95, 123, 123,
## $ drat  <dbl> 3.90, 3.90, 3.85, 3.08, 3.15, 2.76, 3.21, 3.69, 3.92
## $ wt    <dbl> 2.620, 2.875, 2.320, 3.215, 3.440, 3.460, 3.570, 3.1
## $ qsec  <dbl> 16.46, 17.02, 18.61, 19.44, 17.02, 20.22, 15.84, 20.
## $ vs    <dbl> 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1
## $ am    <dbl> 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1
## $ gear  <dbl> 4, 4, 4, 3, 3, 3, 3, 4, 4, 4, 4, 3, 3, 3, 3, 3, 3, 4
## $ carb  <dbl> 4, 4, 1, 1, 2, 1, 4, 2, 2, 4, 4, 3, 3, 3, 4, 4, 4, 1
```

Filter observations

- The function `filter()` is used to subset data with matching logical conditions.
- It is Similar to `base::which()` or `subset()`

```
str(dfm);
```

```
## tibble [32 x 12] (S3: tbl_df/tbl/data.frame)
## $ car : chr [1:32] "Mazda RX4" "Mazda RX4 Wag" "Datsun 710" "Hon
## $ mpg : num [1:32] 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2
## $ cyl : num [1:32] 6 6 4 6 8 6 8 4 4 6 ...
## $ disp: num [1:32] 160 160 108 258 360 ...
## $ hp : num [1:32] 110 110 93 110 175 105 245 62 95 123 ...
## $ drat: num [1:32] 3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.
## $ wt : num [1:32] 2.62 2.88 2.32 3.21 3.44 ...
## $ qsec: num [1:32] 16.5 17 18.6 19.4 17 ...
## $ vs : num [1:32] 0 0 1 1 0 1 0 1 1 1 ...
## $ am : num [1:32] 1 1 1 0 0 0 0 0 0 0 ...
## $ gear: num [1:32] 4 4 4 3 3 3 3 4 4 4 ...
## $ carb: num [1:32] 4 4 1 1 2 1 4 2 2 4 ...
```

Filter observations

- The data set is data frame with 32 observations on 11 (numeric) variables.
 - ▶ [, 1] mpg Miles/(US) gallon
 - ▶ [, 2] cyl Number of cylinders
 - ▶ [, 3] disp Displacement (cu.in.)
 - ▶ [, 4] hp Gross horsepower
 - ▶ [, 5] drat Rear axle ratio
 - ▶ [, 6] wt Weight (1000 lbs)
 - ▶ [, 7] qsec 1/4 mile time
 - ▶ [, 8] vs Engine (0 = V-shaped, 1 = straight)
 - ▶ [, 9] am Transmission (0 = automatic, 1 = manual)
 - ▶ [,10] gear Number of forward gears
 - ▶ [,11] carb Number of carburetors

Filter observations

- `base::unique()` returns unique values of a variable/data frame.

```
unique(dfm$cyl);
```

```
## [1] 6 4 8
```

```
unique(dfm$disp);
```

```
## [1] 160.0 108.0 258.0 360.0 225.0 146.7 140.8 167.6 275.8 472.0
```

```
## [13] 78.7 75.7 71.1 120.1 318.0 304.0 350.0 400.0 79.0 120.3
```

```
## [25] 145.0 301.0 121.0
```

```
unique(dfm$gear);
```

```
## [1] 4 3 5
```

Filter observations

- The command in dplyr for subsetting rows is `filter`. Try `?filter`
- Note, no `$` or subsetting is necessary. R “knows” `mpg` refers to a column of `dfm`.

```
library(dplyr);
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

```
filter(dfm, mpg > 20);
```

```
## # A tibble: 14 x 12
```

```
##      car      mpg    cyl  disp    hp  drat    wt  qsec    vs
```

```
##      <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
```

Filter observations

- Multiple arguments can be provided to `filter()`.

```
filter(dfm, mpg > 20, cyl == 4);
```

```
## # A tibble: 11 x 12
```

##	car	mpg	cyl	disp	hp	drat	wt	qsec	vs	<dbl>
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
##	1 Datsun 710	22.8	4	108	93	3.85	2.32	18.6	1	
##	2 Merc 240D	24.4	4	147.	62	3.69	3.19	20	1	
##	3 Merc 230	22.8	4	141.	95	3.92	3.15	22.9	1	
##	4 Fiat 128	32.4	4	78.7	66	4.08	2.2	19.5	1	
##	5 Honda Civic	30.4	4	75.7	52	4.93	1.62	18.5	1	
##	6 Toyota Cor~	33.9	4	71.1	65	4.22	1.84	19.9	1	
##	7 Toyota Cor~	21.5	4	120.	97	3.7	2.46	20.0	1	
##	8 Fiat X1-9	27.3	4	79	66	4.08	1.94	18.9	1	
##	9 Porsche 91~	26	4	120.	91	4.43	2.14	16.7	0	
##	10 Lotus Euro~	30.4	4	95.1	113	3.77	1.51	16.9	1	
##	11 Volvo 142E	21.4	4	121	109	4.11	2.78	18.6	1	

Filter observations

You can have multiple logical conditions using the following:

- `&` : AND
- `|` : OR

By default, you can separate conditions by commas, and `filter` assumes these statements are joined by `&`:

```
filter(dfm, mpg > 20 & cyl == 4);
```

```
## # A tibble: 11 x 12
```

##	car	mpg	cyl	disp	hp	drat	wt	qsec	vs	<dbl>
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
##	1 Datsun 710	22.8	4	108	93	3.85	2.32	18.6	1	
##	2 Merc 240D	24.4	4	147.	62	3.69	3.19	20	1	
##	3 Merc 230	22.8	4	141.	95	3.92	3.15	22.9	1	
##	4 Fiat 128	32.4	4	78.7	66	4.08	2.2	19.5	1	
##	5 Honda Civic	30.4	4	75.7	52	4.93	1.62	18.5	1	
##	6 Toyota Cor~	33.9	4	71.1	65	4.22	1.84	19.9	1	
##	7 Toyota Cor~	21.5	4	120.	97	3.7	2.46	20.0	1	
##	8 Fiat X1-9	27.3	4	79	66	4.08	1.94	18.9	1	
##	9 Porsche 91~	26	4	120.	91	4.43	2.14	16.7	0	

Filter observations

If you want OR statements, you need to use `|` explicitly:

```
filter(dfm, mpg > 20 | cyl == 4);
```

```
## # A tibble: 14 x 12
```

##	car	mpg	cyl	disp	hp	drat	wt	qsec	vs	<dbl>
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
##	1 Mazda RX4	21	6	160	110	3.9	2.62	16.5	0	
##	2 Mazda RX4 ~	21	6	160	110	3.9	2.88	17.0	0	
##	3 Datsun 710	22.8	4	108	93	3.85	2.32	18.6	1	
##	4 Hornet 4 D~	21.4	6	258	110	3.08	3.22	19.4	1	
##	5 Merc 240D	24.4	4	147.	62	3.69	3.19	20	1	
##	6 Merc 230	22.8	4	141.	95	3.92	3.15	22.9	1	
##	7 Fiat 128	32.4	4	78.7	66	4.08	2.2	19.5	1	
##	8 Honda Civic	30.4	4	75.7	52	4.93	1.62	18.5	1	
##	9 Toyota Cor~	33.9	4	71.1	65	4.22	1.84	19.9	1	
##	10 Toyota Cor~	21.5	4	120.	97	3.7	2.46	20.0	1	
##	11 Fiat X1-9	27.3	4	79	66	4.08	1.94	18.9	1	
##	12 Porsche 91~	26	4	120.	91	4.43	2.14	16.7	0	
##	13 Lotus Euro~	30.4	4	95.1	113	3.77	1.51	16.9	1	
##	14 Volvo 142E	21.4	4	121	109	4.11	2.78	18.6	1	

Filter observations

- `filter()` is similar to `subset()`.
- `subset()` can select both variables (columns) and observations(rows).
- `filter()` works exclusively for observations(rows).
- Recall: Two useful comparison operators `%in%` and `is.na()`:
 - ▶ `x %in% c("a","b","c")` is used to check if `x` is one of the values in the right hand side.
 - ▶ `is.na()` is used to check if a value is missing.

Filter observations - Removing missing values

- like the function `na.omit()` and `questionr::na.rm()`, the function `tidyr::drop_na()` drop rows containing missing values.

```
exam = read.csv("../data/exam.csv", header=TRUE, sep=",");  
any(is.na(exam));
```

```
## [1] TRUE
```

```
which(is.na(exam), arr.ind=TRUE);
```

```
##      row col  
## [1,]    5   4
```

```
library(tidyr);  
exam%>%drop_na();
```

```
##      ID    T1    T2    T3  
## 1     1 21.5 16.5 23.5  
## 2     2 21.0 16.0 20.0  
## 3     3 30.5 31.0 30.0  
## 4     4 23.0 18.5 25.5  
## 5     6 27.5 27.0 35.5
```

Select/drop variables

Recall that we can grab the carb column using the \$ operator.

```
dfm$carb;
```

```
##      [1] 4 4 1 1 2 1 4 2 2 4 4 3 3 3 4 4 4 1 2 1 1 2 2 4 2 1 2 2 4 6
```

Select/drop variables

- `select()` allows you to select specific **columns**(variables) from your data.

```
select(dfm, car, mpg);
```

```
## # A tibble: 32 x 2
##   car                mpg
##   <chr>             <dbl>
## 1 Mazda RX4         21
## 2 Mazda RX4 Wag     21
## 3 Datsun 710        22.8
## 4 Hornet 4 Drive    21.4
## 5 Hornet Sportabout 18.7
## 6 Valiant           18.1
## 7 Duster 360        14.3
## 8 Merc 240D         24.4
## 9 Merc 230          22.8
## 10 Merc 280         19.2
## # ... with 22 more rows
```

Select/drop variables

The select command from dplyr allows you to subset columns of

```
select(dfm, car, mpg, cyl);
```

```
## # A tibble: 32 x 3
```

```
##   car          mpg   cyl
##   <chr>      <dbl> <dbl>
## 1 Mazda RX4      21     6
## 2 Mazda RX4 Wag  21     6
## 3 Datsun 710     22.8    4
## 4 Hornet 4 Drive  21.4    6
## 5 Hornet Sportabout 18.7    8
## 6 Valiant        18.1    6
## 7 Duster 360     14.3    8
## 8 Merc 240D      24.4    4
## 9 Merc 230       22.8    4
## 10 Merc 280      19.2    6
## # ... with 22 more rows
```

```
select(dfm, starts_with("c"));
```

```
## # A tibble: 32 x 3
```

```
##   car          cyl carb
##   <chr>      <dbl> <dbl>
```

See the Select “helpers”

Run the command:

```
library(tidyselect);  
?tidyselect::select_helpers
```

Here are a few:

```
any_of()  
all_of()  
last_col()  
starts_with()  
ends_with()  
contains() # like searching  
matches() # Matches a regular expression - cover later
```

Select/drop variables

- `select()` allows you to **drop** specific columns from your data.

```
select(dfm, -c(vs, am, gear, carb));
```

```
## # A tibble: 32 x 8
```

##	car	mpg	cyl	disp	hp	drat	wt	qsec
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
##	1 Mazda RX4	21	6	160	110	3.9	2.62	16.5
##	2 Mazda RX4 Wag	21	6	160	110	3.9	2.88	17.0
##	3 Datsun 710	22.8	4	108	93	3.85	2.32	18.6
##	4 Hornet 4 Drive	21.4	6	258	110	3.08	3.22	19.4
##	5 Hornet Sportabout	18.7	8	360	175	3.15	3.44	17.0
##	6 Valiant	18.1	6	225	105	2.76	3.46	20.2
##	7 Duster 360	14.3	8	360	245	3.21	3.57	15.8
##	8 Merc 240D	24.4	4	147.	62	3.69	3.19	20
##	9 Merc 230	22.8	4	141.	95	3.92	3.15	22.9
##	10 Merc 280	19.2	6	168.	123	3.92	3.44	18.3
##	# ... with 22 more rows							

Combining filter and select

You can combine filter and select to subset the rows and columns, respectively, of a data.frame:

```
select(filter(dfm, mpg > 20, cyl == 4), car, cyl, hp);
```

```
## # A tibble: 11 x 3
```

```
##      car      cyl    hp
##    <chr>    <dbl> <dbl>
## 1 Datsun 710      4     93
## 2 Merc 240D      4     62
## 3 Merc 230       4     95
## 4 Fiat 128       4     66
## 5 Honda Civic    4     52
## 6 Toyota Corolla 4     65
## 7 Toyota Corona  4     97
## 8 Fiat X1-9      4     66
## 9 Porsche 914-2   4     91
## 10 Lotus Europa   4    113
## 11 Volvo 142E     4    109
```

In R, the common way to perform multiple operations is to wrap functions around each other in a nested way such as above.

The %>% operator

Recently, the pipe %>% makes things such as this much more readable. It reads left side “pipes” into right side. RStudio CMD/Ctrl + Shift + M shortcut. Pipe dfm into filter, then pipe that into select:

```
dfm %>% filter(mpg > 20, cyl == 4) %>% select(car,cyl, hp);
```

```
## # A tibble: 11 x 3
##   car          cyl    hp
##   <chr>        <dbl> <dbl>
## 1 Datsun 710      4     93
## 2 Merc 240D       4     62
## 3 Merc 230        4     95
## 4 Fiat 128        4     66
## 5 Honda Civic     4     52
## 6 Toyota Corolla  4     65
## 7 Toyota Corona   4     97
## 8 Fiat X1-9        4     66
## 9 Porsche 914-2    4     91
## 10 Lotus Europa    4    113
## 11 Volvo 142E      4    109
```

The %>% operator

- Powerful trick for coding a sequence of operations
- Output of old operation as the first argument of new operation
- Especially useful in combined with package `ggplot2`

# input	+-----+		+-----+		+-----+		result
# data	%>%	verb	%>%	verb	%>%	verb	-> data
# frame	+-----+		+-----+		+-----+		frame

Add/remove variables

- Adding new columns to a data.frame: base R

You can add a new column, called `newcol` to `dfm`, using the `$` operator:

```
dfm$newcol = dfm$wt/2;  
head(dfm,3)
```

```
## # A tibble: 3 x 13
```

```
##   car      mpg   cyl  disp    hp  drat    wt   qsec    vs    am  ge  
##   <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 Mazd~   21     6   160   110   3.9   2.62  16.5     0     1  
## 2 Mazd~   21     6   160   110   3.9   2.88  17.0     0     1  
## 3 Dats~  22.8    4   108    93   3.85  2.32  18.6     1     1
```

Add/remove variables

- The \$ method is very common.
- The mutate function in dplyr allows you to add variables to your dataset.

```
dfm = mutate(dfm, newcol = wt/2);
```

Creating conditional variables

- One frequently-used tool is creating variables with conditions.
- A general function for creating new variables based on existing variables is the `ifelse()` function, which “returns a value with the same shape as test which is filled with elements selected from either yes or no depending on whether the element of test is TRUE or FALSE.” It is much faster than combining “if(){} else{ }” and “for” loop.

```
ifelse(test, yes, no)
```

```
# test: an object which can be coerced  
#       to logical mode.
```

```
# yes: return values for true elements of test.
```

```
# no: return values for false elements of test.
```

Add/remove variables

Combined with `ifelse(condition, TRUE, FALSE)`, it can give you:

```
dfm = mutate(dfm,
              disp_cat = ifelse(
                disp <= 200,
                "Low",
                ifelse(disp <= 400,
                      "Medium",
                      "High")
                )
              );
head(dfm$disp_cat);

## [1] "Low"    "Low"    "Low"    "Medium" "Medium" "Medium"
```

Add/remove variables

Alternatively, `dplyr::case_when` provides a clean syntax as well.

```
dfm = mutate(dfm,  
  disp_cat2 = case_when(  
    disp <= 200 ~ "Low",  
    disp > 200 & disp <= 400 ~ "Medium",  
    disp > 400 ~ "High")  
);  
head(dfm$disp_cat2);
```

```
## [1] "Low"    "Low"    "Low"    "Medium" "Medium" "Medium"
```


Transmutation

The `transmute` function in `dplyr` combines both the `mutate` and `select` functions. One can create new columns and keep the only the columns wanted:

```
transmute(dfm, newcol2 = wt/2, mpg, hp);
```

```
## # A tibble: 32 x 3
##   newcol2    mpg    hp
##   <dbl> <dbl> <dbl>
## 1     1.31    21    110
## 2     1.44    21    110
## 3     1.16   22.8     93
## 4     1.61   21.4    110
## 5     1.72   18.7    175
## 6     1.73   18.1    105
## 7     1.78   14.3   245
## 8     1.60   24.4     62
## 9     1.58   22.8     95
## 10    1.72   19.2    123
## # ... with 22 more rows
```

Add/remove variables

- Recall that we can remove a column by assigning to NULL:

```
dfm$newcol = NULL;
```

- The NULL method is still very common.
- Or we can use the select function a minus (-)

```
select(dfm, -newcol);
```

```
## # A tibble: 32 x 14
```

##	car	mpg	cyl	disp	hp	drat	wt	qsec	vs	<
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<
##	1 Mazda RX4	21	6	160	110	3.9	2.62	16.5	0	
##	2 Mazda RX4 ~	21	6	160	110	3.9	2.88	17.0	0	
##	3 Datsun 710	22.8	4	108	93	3.85	2.32	18.6	1	
##	4 Hornet 4 D~	21.4	6	258	110	3.08	3.22	19.4	1	
##	5 Hornet Spo~	18.7	8	360	175	3.15	3.44	17.0	0	
##	6 Valiant	18.1	6	225	105	2.76	3.46	20.2	1	
##	7 Duster 360	14.3	8	360	245	3.21	3.57	15.8	0	
##	8 Merc 240D	24.4	4	147.	62	3.69	3.19	20	1	
##	9 Merc 230	22.8	4	141.	95	3.92	3.15	22.9	1	

Rename variables

We can use the `colnames` function to directly reassign column names of `df`:

```
colnames(dfm)[2:4] = c("MPG", "CYL", "DISP");  
head(dfm);
```

```
## # A tibble: 6 x 15
```

```
##   car      MPG   CYL  DISP    hp  drat    wt  qsec    vs    am  ge  
##   <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 Mazd~    21     6   160   110  3.9   2.62  16.5     0     1  
## 2 Mazd~    21     6   160   110  3.9   2.88  17.0     0     1  
## 3 Dats~   22.8     4   108    93  3.85  2.32  18.6     1     1  
## 4 Horn~   21.4     6   258   110  3.08  3.22  19.4     1     0  
## 5 Horn~   18.7     8   360   175  3.15  3.44  17.0     0     0  
## 6 Vali~   18.1     6   225   105  2.76  3.46  20.2     1     0  
## # ... with 2 more variables: disp_cat <chr>, disp_cat2 <chr>
```

```
colnames(dfm)[2:4] = c("mpg", "cyl", "disp"); #reset
```

Rename variables

- `rename()` function can be used to change variable name.

```
rename(data, new_name = old_name)
```

```
dfm=rename(dfm, WT= wt);  
head(dfm);
```

```
## # A tibble: 6 x 15
```

```
##   car      mpg   cyl  disp    hp  drat    WT   qsec    vs    am  ge  
##   <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>  
## 1 Mazd~    21     6   160   110   3.9   2.62  16.5     0     1  
## 2 Mazd~    21     6   160   110   3.9   2.88  17.0     0     1  
## 3 Dats~   22.8     4   108    93   3.85   2.32  18.6     1     1  
## 4 Horn~   21.4     6   258   110   3.08   3.22  19.4     1     0  
## 5 Horn~   18.7     8   360   175   3.15   3.44  17.0     0     0  
## 6 Vali~   18.1     6   225   105   2.76   3.46  20.2     1     0  
## # ... with 2 more variables: disp_cat <chr>, disp_cat2 <chr>
```

```
dfm=rename(dfm, wt= WT);  #reset
```

Sort data

- The `select` function can reorder columns. Put `newcol` first, then select the rest of columns.

```
select(dfm, cyl, everything());
```

```
## # A tibble: 32 x 15
```

```
##       cyl car           mpg  disp    hp  drat    wt  qsec    vs
##    <dbl> <chr>         <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1      6 Mazda RX4      21    160   110   3.9   2.62  16.5     0
## 2      6 Mazda RX4 ~    21    160   110   3.9   2.88  17.0     0
## 3      4 Datsun 710     22.8   108    93   3.85   2.32  18.6     1
## 4      6 Hornet 4 D~    21.4   258   110   3.08   3.22  19.4     1
## 5      8 Hornet Spo~    18.7   360   175   3.15   3.44  17.0     0
## 6      6 Valiant        18.1   225   105   2.76   3.46  20.2     1
## 7      8 Duster 360     14.3   360   245   3.21   3.57  15.8     0
## 8      4 Merc 240D      24.4   147    62   3.69   3.19   20     1
## 9      4 Merc 230       22.8   141    95   3.92   3.15  22.9     1
## 10     6 Merc 280       19.2   168   123   3.92   3.44  18.3     1
## # ... with 22 more rows, and 3 more variables: newcol <dbl>, disp
## #   disp_cat2 <chr>
```

Sort data

- `arrange()` sorts your data (by rows). In base R, this is commonly done with `order()`.
- By default, `arrange` orders in ascending order.
- Use the `desc` to arrange the rows in descending order:

```
data=data.frame(x = c(2,3,5,1,4), y=c(10, 9, 7, 8, 6));  
arrange(data,x,y);
```

```
##      x  y  
## 1 1  8  
## 2 2 10  
## 3 3  9  
## 4 4  6  
## 5 5  7
```

```
arrange(data, desc(x));
```

```
##      x  y  
## 1 5  7  
## 2 4  6  
## 3 3  9  
## 4 2 10  
## 5 1  8
```

Sort data

- It is a bit more straightforward to mix increasing and decreasing orderings

```
arrange(dfm, mpg, desc(hp));
```

```
## # A tibble: 32 x 15
```

```
##   car                mpg   cyl  disp    hp  drat    wt   qsec    vs
##   <chr>            <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Lincoln Co~    10.4     8  460    215   3     5.42  17.8     0
## 2 Cadillac F~    10.4     8  472    205  2.93   5.25  18.0     0
## 3 Camaro Z28     13.3     8  350    245  3.73   3.84  15.4     0
## 4 Duster 360     14.3     8  360    245  3.21   3.57  15.8     0
## 5 Chrysler I~    14.7     8  440    230  3.23   5.34  17.4     0
## 6 Maserati B~    15       8  301    335  3.54   3.57  14.6     0
## 7 Merc 450SLC    15.2     8  276.    180  3.07   3.78   18      0
## 8 AMC Javelin    15.2     8  304    150  3.15   3.44  17.3     0
## 9 Dodge Chal~    15.5     8  318    150  2.76   3.52  16.9     0
## 10 Ford Pante~    15.8     8  351    264  4.22   3.17  14.5     0
## # ... with 22 more rows, and 3 more variables: newcol <dbl>, disp
## #   disp_cat2 <chr>
```

Remove duplicate rows

- The `distinct()` function is used to eliminate duplicates.
- Remove duplicate rows based on all variables

```
dim(dfm);
```

```
## [1] 32 15
```

```
dfm1 = distinct(dfm);
```

```
dim(dfm1);
```

```
## [1] 32 15
```

- Remove duplicate rows **based on some variables**
 - ▶ `.keep_all` option is used to retain all other variables in the output data frame. It is `FALSE` by default

```
dfm2 = distinct(dfm, mpg, cyl, .keep_all= TRUE)
```

```
dim(dfm2)
```

```
## [1] 27 15
```


License



This work is licensed under a [Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License](#).